

$o(n)$ -coloring 3-colorable graphs in $O(n)$ time

Tom Tseng

In this writeup, we give an algorithm that takes as input a 3-colorable graph in adjacency matrix format and outputs an $O(n/\log \log n)$ coloring in $O(n)$ time under the word RAM model.

1 Notation

We refer to the number of vertices of the input graph as n . We'll let $w \geq \log_2 n$ denote the word size. We'll assume $w \in \Theta(\log n)$, because if we had $w \in \omega(\log n)$, we could just ignore the higher bits of every word.

To more closely follow the notation used in many programming languages for bitwise logical operators, we'll use $\&$ to denote bitwise conjunction (AND), $|$ to denote bitwise disjunction (OR), and \sim to denote bitwise negation (NOT). More specifically, if we have two boolean vectors v and w of length ℓ , then the results of $v \& w$, $v | w$, and $\sim v$ are all boolean vectors of length ℓ such that

$$(v \& w)_i = v_i \wedge w_i, \quad (v | w)_i = v_i \vee w_i, \quad (\sim v)_i = \neg v_i.$$

When A is a matrix and v is a vector, $A \cdot v$ represents boolean matrix multiplication, that is,

$$(A \cdot v)_i = \bigvee_j A_{i,j} \wedge v_j.$$

We'll use “1” and “true” interchangeably and “0” and “false” interchangeably.

2 Sketch of the algorithm

The high-level idea is as follows: we break the vertices into chunks of size k , and for the subgraph induced by each chunk, we 3-color the subgraph by brute force (using a different set of 3 colors for each subgraph). This gives a $3n/k$ coloring of the graph. We set k to be just a little bigger than a constant. To keep the runtime linear, we use some bit hackery to get some word-level parallelism.

Let the input be given in adjacency matrix format. Pick $k \in \Omega(\log \log n)$ such that $3^k(k+1) \leq w$, e.g. $k = \log_4 w$ will suffice. We partition the vertices into n/k contiguous chunks of k , and if we can 3-color the subgraph induced by each chunk in $O(k)$ time, then we'll achieve our desired result.

We can represent a 3-coloring of a graph of k vertices by three k -length bit vectors. The j -th bit of the i -th vector is set if and only if the j -th vertex has color i . The idea here is that if we have the three k -length bit vectors v_0, v_1, v_2 representing a 3-coloring as well as the adjacency matrix A of the k -vertex graph, we can check that the coloring is valid by checking that $(A \cdot v_i) \& v_i = \vec{0}$ for each i . This is because the j -th bit of $A \cdot v_i$ is set if the j -th vertex has any neighbors of color i , so then ANDing with v_i tells us about which i -colored vertices have i -colored neighbors. Due to how small k is, we can check all possible 3-colorings for validity in parallel.

We start by precomputing some constants. Because k is so small, we can pack a the aforementioned representation of all 3^k colorings into three words u_0, u_1, u_2 . Each word u_i is broken into 3^k blocks where each block is $(k+1)$ bits wide. The k -length bit vector for the i -th color of the j -th possible 3-coloring is right-aligned within the j -th block of u_i . We also precompute B_L to be a word broken into the same 3^k blocks where each block holds just a single left-aligned 1 bit, and precompute B_R to be a word broken in 3^k blocks where each block holds just a single right-aligned 1 bit.

Iterate over each chunk of k vertices and perform the following: consider the subgraph induced by the k vertices. We'll perform the parallel boolean matrix multiplication now. For each $r = 0, 1, \dots, k-1$, we fetch the r -th row of the $k \times k$ adjacency matrix in constant time by jumping to the appropriate place in the input and doing some bit masking and shifting. Multiply the word by B_R so that we now have a word w_r consisting of 3^k copies of row r of the adjacency matrix. Now $w_r \& u_i$ is a word of 3^k blocks where the j -th block is non-zero iff the r -th entry of the corresponding boolean matrix product is non-zero. Then $z_{r,i} = (\sim(B_L - (w_r \& u_i))) \& B_L$ is a word of 3^k blocks where the j -th block has its high-order bit set iff the r -th entry of the corresponding boolean matrix product is non-zero. (Computing each $z_{r,i}$ is constant time, so computing all of them takes $O(k)$ time.) Shift and OR each $z_{r,i}$ appropriately to get words y_i of 3^k blocks where the j -th block has the result of the boolean matrix product corresponding to color i of the j -th coloring. Compute $y = (y_0 \& u_0) \mid (y_1 \& u_1) \mid (y_2 \& u_2)$, which has that the j -th block is all zeroes if the j -th coloring is valid. Compute $x = (B_L - y) \& B_L$, which has that its j -th block has its high-order bit set to 1 if the j -th coloring is valid. Binary search for a set bit in x in $O(\log w) = O(k)$ using lots of masking, and after finding that bit, we can read off a 3-coloring for the subgraph by indexing appropriately into u_0, u_1, u_2 . This is all $O(k)$ time for a chunk of k vertices.

We do this for n/k chunks of k vertices, so this takes $n/k \cdot O(k) = O(n)$ time. The number of colors used over the whole graph is $3n/k = O(n/\log \log n) \subseteq o(n)$.